

# A4-Q3: Convolutional Autoencoders

## Preliminaries

```
In [70]: import numpy as np
import torch
import torch.nn as nn
import torchvision
import matplotlib.pyplot as plt
from tqdm import tqdm
```

```
In [ ]:
```

```
In [71]: # In case you are fortunate enough to have access to a GPU...
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

## Dataset: MNIST

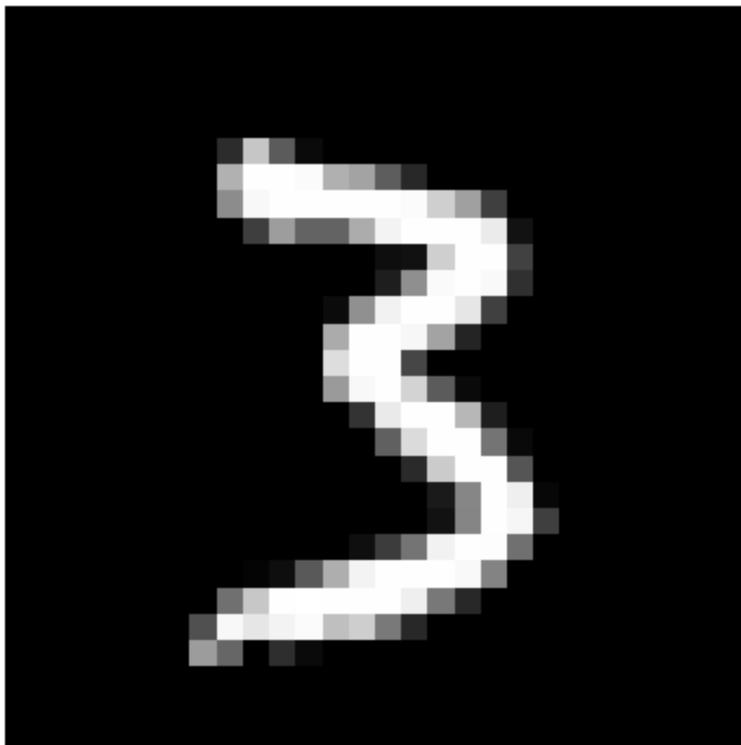
```
In [72]: # You can change img_size to 14 if you want to use smaller (14x14) images.
img_size = 28
ds_full = torchvision.datasets.MNIST('./files/', train=True, download=True,
                                      transform=torchvision.transforms.Compose([
                                          torchvision.transforms.Resize((img_size, img_size)),
                                          torchvision.transforms.ToTensor(),
                                      ]))
```

```
In [73]: ds = torch.utils.data.Subset(ds_full, range(1024))
```

```
In [74]: def Draw(x):
    with torch.no_grad():
        plt.imshow(x.squeeze().detach().numpy(), cmap='gray');
        plt.axis('off');
```

```
In [77]: with torch.no_grad():
    x,t = ds.__getitem__(130)
    Draw(x)
    plt.title(f'Size: {list(x.size())}')
plt.show()
```

Size: [1, 28, 28]



## Create some `DataLoader`s

```
In [78]: # Batched, for training  
batch_size_train = 8  
train_dl = torch.utils.data.DataLoader(ds, batch_size=batch_size_train, shuffle=True)
```

```
In [79]: # A single batch, for plotting  
train_all = torch.utils.data.DataLoader(ds, batch_size=1024, shuffle=False)
```

## A. Complete the `ConvAE` class

```
In [80]: class ConvAE(nn.Module):  
    ...  
    net = ConvAE(img_size=28, embedding_dim=3)  
  
    Create a convolutional autoencoder for input images of size (img_size x img_size)  
    with an embedding (latent) layer of (embedding_dim) neurons.  
  
    Inputs:  
        img_size      size of input images, [1, img_size, img_size]  
        embedding_dim number of nodes in embedding (latent) layer  
  
    Usage:  
        net = ConvAE()  
        y = net(x)  
        h = net.encode(x) # returns latent vectors
```

```

    ...
    def __init__(self, img_size=28, embedding_dim=3):
        self.img_size = img_size
        self.embedding_dim = embedding_dim
        self.losses = []
        super().__init__()

        ##### YOUR CODE HERE #####
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 16, kernel_size=3, stride=2, padding=1),
            nn.ReLU(True),
            nn.Conv2d(16, 32, kernel_size=3, stride=2, padding=1),
            nn.ReLU(True),
            nn.Conv2d(32, 64, kernel_size=3, stride=2, padding=1),
            nn.ReLU(True),
            nn.Flatten()
        )
        self.fc_enc = nn.Sequential(
            nn.Linear(64 * 4 * 4, 128),
            nn.ReLU(True),
            nn.Linear(128, embedding_dim),
            nn.Tanh()
        )

        self.fc_dec = nn.Sequential(
            nn.Linear(embedding_dim, 128),
            nn.ReLU(True),
            nn.Linear(128, 64 * 4 * 4),
            nn.ReLU(True),
            nn.Unflatten(1, (64, 4, 4))
        )
        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(64, 32, kernel_size=3, stride=2, padding=1, output_padding=1),
            nn.ReLU(True),
            nn.ConvTranspose2d(32, 16, kernel_size=3, stride=2, padding=1, output_padding=1),
            nn.ReLU(True),
            nn.ConvTranspose2d(16, 1, kernel_size=3, stride=2, padding=1, output_padding=1),
            nn.Sigmoid()
        )

    def encode(self, x):
        ##### YOUR CODE HERE #####
        x = self.encoder(x)
        h = self.fc_enc(x)
        return h

    def decode(self, x):
        ##### YOUR CODE HERE #####
        h = self.fc_dec(x)
        x = self.decoder(h)
        x = x[:, :, :self.img_size, :self.img_size]
        return x

    def forward(self, x):

```

```
===== YOUR CODE HERE =====
h = self.encode(x)
y = self.decode(h)
return y
```

## B. Create and train the network

```
In [81]: net = ConvAE(img_size=img_size, embedding_dim=3)
```

```
In [82]: # Train it
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3, weight_decay=1e-5)
criterion = nn.MSELoss()
epochs = 100

for epoch in range(epochs):
    net.train()
    running_loss = 0

    for x, _ in tqdm(train_dl, desc=f"Epoch {epoch+1}/{epochs}"):
        x = x.to(device)

        y_pred = net(x)

        loss = criterion(y_pred, x)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item()

    avg_loss = running_loss / len(train_dl)
    net.losses.append(avg_loss)
    print(f"Epoch [{epoch+1}/{epochs}] Average Loss: {avg_loss:.6f}")

print("\nTraining complete!")
```

```
Epoch 1/100: 100%|██████████| 128/128 [00:00<00:00, 298.93it/s]
Epoch [1/100] Average Loss: 0.094311
Epoch 2/100: 100%|██████████| 128/128 [00:00<00:00, 333.11it/s]
Epoch [2/100] Average Loss: 0.066325
Epoch 3/100: 100%|██████████| 128/128 [00:00<00:00, 297.87it/s]
Epoch [3/100] Average Loss: 0.066098
Epoch 4/100: 100%|██████████| 128/128 [00:00<00:00, 305.37it/s]
Epoch [4/100] Average Loss: 0.066066
Epoch 5/100: 100%|██████████| 128/128 [00:00<00:00, 316.04it/s]
Epoch [5/100] Average Loss: 0.066016
Epoch 6/100: 100%|██████████| 128/128 [00:00<00:00, 282.23it/s]
Epoch [6/100] Average Loss: 0.065992
Epoch 7/100: 100%|██████████| 128/128 [00:00<00:00, 321.09it/s]
Epoch [7/100] Average Loss: 0.066038
```

Epoch 8/100: 100%|██████████| 128/128 [00:00<00:00, 326.99it/s]

Epoch [8/100] Average Loss: 0.065589

Epoch 9/100: 100%|██████████| 128/128 [00:00<00:00, 336.16it/s]

Epoch [9/100] Average Loss: 0.061889

Epoch 10/100: 100%|██████████| 128/128 [00:00<00:00, 301.06it/s]

Epoch [10/100] Average Loss: 0.058413

Epoch 11/100: 100%|██████████| 128/128 [00:00<00:00, 319.25it/s]

Epoch [11/100] Average Loss: 0.053229

Epoch 12/100: 100%|██████████| 128/128 [00:00<00:00, 347.75it/s]

Epoch [12/100] Average Loss: 0.050628

Epoch 13/100: 100%|██████████| 128/128 [00:00<00:00, 326.38it/s]

Epoch [13/100] Average Loss: 0.047173

Epoch 14/100: 100%|██████████| 128/128 [00:00<00:00, 314.01it/s]

Epoch [14/100] Average Loss: 0.044147

Epoch 15/100: 100%|██████████| 128/128 [00:00<00:00, 335.90it/s]

Epoch [15/100] Average Loss: 0.042780

Epoch 16/100: 100%|██████████| 128/128 [00:00<00:00, 331.97it/s]

Epoch [16/100] Average Loss: 0.041491

Epoch 17/100: 100%|██████████| 128/128 [00:00<00:00, 334.59it/s]

Epoch [17/100] Average Loss: 0.040398

Epoch 18/100: 100%|██████████| 128/128 [00:00<00:00, 336.09it/s]

Epoch [18/100] Average Loss: 0.039726

Epoch 19/100: 100%|██████████| 128/128 [00:00<00:00, 339.70it/s]

Epoch [19/100] Average Loss: 0.038845

Epoch 20/100: 100%|██████████| 128/128 [00:00<00:00, 312.60it/s]

Epoch [20/100] Average Loss: 0.038110

Epoch 21/100: 100%|██████████| 128/128 [00:00<00:00, 328.88it/s]

Epoch [21/100] Average Loss: 0.037718

Epoch 22/100: 100%|██████████| 128/128 [00:00<00:00, 339.22it/s]

Epoch [22/100] Average Loss: 0.037302

Epoch 23/100: 100%|██████████| 128/128 [00:00<00:00, 336.29it/s]

Epoch [23/100] Average Loss: 0.036799

Epoch 24/100: 100%|██████████| 128/128 [00:00<00:00, 336.74it/s]

Epoch [24/100] Average Loss: 0.036237

Epoch 25/100: 100%|██████████| 128/128 [00:00<00:00, 337.17it/s]

Epoch [25/100] Average Loss: 0.036009

Epoch 26/100: 100%|██████████| 128/128 [00:00<00:00, 334.59it/s]

Epoch [26/100] Average Loss: 0.035396

Epoch 27/100: 100%|██████████| 128/128 [00:00<00:00, 330.67it/s]

Epoch [27/100] Average Loss: 0.035097

Epoch 28/100: 100%|██████████| 128/128 [00:00<00:00, 346.76it/s]

Epoch [28/100] Average Loss: 0.034641

Epoch 29/100: 100%|██████████| 128/128 [00:00<00:00, 342.58it/s]

Epoch [29/100] Average Loss: 0.034551

Epoch 30/100: 100%|██████████| 128/128 [00:00<00:00, 336.88it/s]

Epoch [30/100] Average Loss: 0.034383

Epoch 31/100: 100%|██████████| 128/128 [00:00<00:00, 352.53it/s]

Epoch [31/100] Average Loss: 0.033887

Epoch 32/100: 100%|██████████| 128/128 [00:00<00:00, 353.96it/s]

Epoch [32/100] Average Loss: 0.033738  
Epoch 33/100: 100% |██████████| 128/128 [00:00<00:00, 356.37it/s]  
Epoch [33/100] Average Loss: 0.033376  
Epoch 34/100: 100% |██████████| 128/128 [00:00<00:00, 314.48it/s]  
Epoch [34/100] Average Loss: 0.033217  
Epoch 35/100: 100% |██████████| 128/128 [00:00<00:00, 325.81it/s]  
Epoch [35/100] Average Loss: 0.033032  
Epoch 36/100: 100% |██████████| 128/128 [00:00<00:00, 319.99it/s]  
Epoch [36/100] Average Loss: 0.032820  
Epoch 37/100: 100% |██████████| 128/128 [00:00<00:00, 334.34it/s]  
Epoch [37/100] Average Loss: 0.032732  
Epoch 38/100: 100% |██████████| 128/128 [00:00<00:00, 344.01it/s]  
Epoch [38/100] Average Loss: 0.032297  
Epoch 39/100: 100% |██████████| 128/128 [00:00<00:00, 355.82it/s]  
Epoch [39/100] Average Loss: 0.032170  
Epoch 40/100: 100% |██████████| 128/128 [00:00<00:00, 329.93it/s]  
Epoch [40/100] Average Loss: 0.032210  
Epoch 41/100: 100% |██████████| 128/128 [00:00<00:00, 344.81it/s]  
Epoch [41/100] Average Loss: 0.031973  
Epoch 42/100: 100% |██████████| 128/128 [00:00<00:00, 301.96it/s]  
Epoch [42/100] Average Loss: 0.031800  
Epoch 43/100: 100% |██████████| 128/128 [00:00<00:00, 310.39it/s]  
Epoch [43/100] Average Loss: 0.031709  
Epoch 44/100: 100% |██████████| 128/128 [00:00<00:00, 283.62it/s]  
Epoch [44/100] Average Loss: 0.031687  
Epoch 45/100: 100% |██████████| 128/128 [00:00<00:00, 310.97it/s]  
Epoch [45/100] Average Loss: 0.031582  
Epoch 46/100: 100% |██████████| 128/128 [00:00<00:00, 293.78it/s]  
Epoch [46/100] Average Loss: 0.031273  
Epoch 47/100: 100% |██████████| 128/128 [00:00<00:00, 335.98it/s]  
Epoch [47/100] Average Loss: 0.031180  
Epoch 48/100: 100% |██████████| 128/128 [00:00<00:00, 300.94it/s]  
Epoch [48/100] Average Loss: 0.031075  
Epoch 49/100: 100% |██████████| 128/128 [00:00<00:00, 323.50it/s]  
Epoch [49/100] Average Loss: 0.030884  
Epoch 50/100: 100% |██████████| 128/128 [00:00<00:00, 336.76it/s]  
Epoch [50/100] Average Loss: 0.030669  
Epoch 51/100: 100% |██████████| 128/128 [00:00<00:00, 332.24it/s]  
Epoch [51/100] Average Loss: 0.030905  
Epoch 52/100: 100% |██████████| 128/128 [00:00<00:00, 326.23it/s]  
Epoch [52/100] Average Loss: 0.030613  
Epoch 53/100: 100% |██████████| 128/128 [00:00<00:00, 329.22it/s]  
Epoch [53/100] Average Loss: 0.030437  
Epoch 54/100: 100% |██████████| 128/128 [00:00<00:00, 327.41it/s]  
Epoch [54/100] Average Loss: 0.030602  
Epoch 55/100: 100% |██████████| 128/128 [00:00<00:00, 334.88it/s]  
Epoch [55/100] Average Loss: 0.030307  
Epoch 56/100: 100% |██████████| 128/128 [00:00<00:00, 330.00it/s]  
Epoch [56/100] Average Loss: 0.030160

Epoch 57/100: 100% |██████████| 128/128 [00:00<00:00, 332.22it/s]  
Epoch [57/100] Average Loss: 0.030249

Epoch 58/100: 100% |██████████| 128/128 [00:00<00:00, 322.36it/s]  
Epoch [58/100] Average Loss: 0.030046

Epoch 59/100: 100% |██████████| 128/128 [00:00<00:00, 336.16it/s]  
Epoch [59/100] Average Loss: 0.029871

Epoch 60/100: 100% |██████████| 128/128 [00:00<00:00, 329.00it/s]  
Epoch [60/100] Average Loss: 0.029734

Epoch 61/100: 100% |██████████| 128/128 [00:00<00:00, 334.79it/s]  
Epoch [61/100] Average Loss: 0.029833

Epoch 62/100: 100% |██████████| 128/128 [00:00<00:00, 276.68it/s]  
Epoch [62/100] Average Loss: 0.029598

Epoch 63/100: 100% |██████████| 128/128 [00:00<00:00, 321.75it/s]  
Epoch [63/100] Average Loss: 0.029699

Epoch 64/100: 100% |██████████| 128/128 [00:00<00:00, 318.65it/s]  
Epoch [64/100] Average Loss: 0.029592

Epoch 65/100: 100% |██████████| 128/128 [00:00<00:00, 291.10it/s]  
Epoch [65/100] Average Loss: 0.029303

Epoch 66/100: 100% |██████████| 128/128 [00:00<00:00, 336.37it/s]  
Epoch [66/100] Average Loss: 0.029678

Epoch 67/100: 100% |██████████| 128/128 [00:00<00:00, 220.99it/s]  
Epoch [67/100] Average Loss: 0.029353

Epoch 68/100: 100% |██████████| 128/128 [00:00<00:00, 222.24it/s]  
Epoch [68/100] Average Loss: 0.029115

Epoch 69/100: 100% |██████████| 128/128 [00:00<00:00, 275.35it/s]  
Epoch [69/100] Average Loss: 0.029124

Epoch 70/100: 100% |██████████| 128/128 [00:00<00:00, 293.04it/s]  
Epoch [70/100] Average Loss: 0.028942

Epoch 71/100: 100% |██████████| 128/128 [00:00<00:00, 302.34it/s]  
Epoch [71/100] Average Loss: 0.028969

Epoch 72/100: 100% |██████████| 128/128 [00:00<00:00, 273.53it/s]  
Epoch [72/100] Average Loss: 0.028793

Epoch 73/100: 100% |██████████| 128/128 [00:00<00:00, 297.84it/s]  
Epoch [73/100] Average Loss: 0.028687

Epoch 74/100: 100% |██████████| 128/128 [00:00<00:00, 299.20it/s]  
Epoch [74/100] Average Loss: 0.029036

Epoch 75/100: 100% |██████████| 128/128 [00:00<00:00, 284.33it/s]  
Epoch [75/100] Average Loss: 0.028865

Epoch 76/100: 100% |██████████| 128/128 [00:00<00:00, 314.80it/s]  
Epoch [76/100] Average Loss: 0.028479

Epoch 77/100: 100% |██████████| 128/128 [00:00<00:00, 311.05it/s]  
Epoch [77/100] Average Loss: 0.028582

Epoch 78/100: 100% |██████████| 128/128 [00:00<00:00, 304.87it/s]  
Epoch [78/100] Average Loss: 0.028654

Epoch 79/100: 100% |██████████| 128/128 [00:00<00:00, 314.92it/s]  
Epoch [79/100] Average Loss: 0.028408

Epoch 80/100: 100% |██████████| 128/128 [00:00<00:00, 271.30it/s]  
Epoch [80/100] Average Loss: 0.028271

Epoch 81/100: 100% |██████████| 128/128 [00:00<00:00, 307.28it/s]

```
Epoch [81/100] Average Loss: 0.028317
Epoch 82/100: 100%|██████████| 128/128 [00:00<00:00, 292.36it/s]
Epoch [82/100] Average Loss: 0.028206
Epoch 83/100: 100%|██████████| 128/128 [00:00<00:00, 311.44it/s]
Epoch [83/100] Average Loss: 0.028313
Epoch 84/100: 100%|██████████| 128/128 [00:00<00:00, 306.14it/s]
Epoch [84/100] Average Loss: 0.028031
Epoch 85/100: 100%|██████████| 128/128 [00:00<00:00, 306.73it/s]
Epoch [85/100] Average Loss: 0.028088
Epoch 86/100: 100%|██████████| 128/128 [00:00<00:00, 302.02it/s]
Epoch [86/100] Average Loss: 0.028096
Epoch 87/100: 100%|██████████| 128/128 [00:00<00:00, 298.97it/s]
Epoch [87/100] Average Loss: 0.027919
Epoch 88/100: 100%|██████████| 128/128 [00:00<00:00, 287.16it/s]
Epoch [88/100] Average Loss: 0.027987
Epoch 89/100: 100%|██████████| 128/128 [00:00<00:00, 308.92it/s]
Epoch [89/100] Average Loss: 0.027994
Epoch 90/100: 100%|██████████| 128/128 [00:00<00:00, 292.34it/s]
Epoch [90/100] Average Loss: 0.027649
Epoch 91/100: 100%|██████████| 128/128 [00:00<00:00, 305.27it/s]
Epoch [91/100] Average Loss: 0.027555
Epoch 92/100: 100%|██████████| 128/128 [00:00<00:00, 312.63it/s]
Epoch [92/100] Average Loss: 0.027720
Epoch 93/100: 100%|██████████| 128/128 [00:00<00:00, 308.72it/s]
Epoch [93/100] Average Loss: 0.027724
Epoch 94/100: 100%|██████████| 128/128 [00:00<00:00, 303.76it/s]
Epoch [94/100] Average Loss: 0.027465
Epoch 95/100: 100%|██████████| 128/128 [00:00<00:00, 309.01it/s]
Epoch [95/100] Average Loss: 0.027553
Epoch 96/100: 100%|██████████| 128/128 [00:00<00:00, 310.87it/s]
Epoch [96/100] Average Loss: 0.027673
Epoch 97/100: 100%|██████████| 128/128 [00:00<00:00, 304.92it/s]
Epoch [97/100] Average Loss: 0.027372
Epoch 98/100: 100%|██████████| 128/128 [00:00<00:00, 310.46it/s]
Epoch [98/100] Average Loss: 0.027316
Epoch 99/100: 100%|██████████| 128/128 [00:00<00:00, 267.35it/s]
Epoch [99/100] Average Loss: 0.027343
Epoch 100/100: 100%|██████████| 128/128 [00:00<00:00, 272.13it/s]
Epoch [100/100] Average Loss: 0.027298
```

Training complete!

```
In [83]: # Uncomment the following line if you want to save your network.
#torch.save(net.to('cpu'), 'my_ConvAE.pt')
# The corresponding code to reload the network is below.
#net = torch.load('my_ConvAE.pt')
# Remember to send it to the GPU, if you're using one.
net.to(device)
```

```

Out[83]: ConvAE(
    (encoder): Sequential(
        (0): Conv2d(1, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (1): ReLU(inplace=True)
        (2): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (3): ReLU(inplace=True)
        (4): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
        (5): ReLU(inplace=True)
        (6): Flatten(start_dim=1, end_dim=-1)
    )
    (fc_enc): Sequential(
        (0): Linear(in_features=1024, out_features=128, bias=True)
        (1): ReLU(inplace=True)
        (2): Linear(in_features=128, out_features=3, bias=True)
        (3): Tanh()
    )
    (fc_dec): Sequential(
        (0): Linear(in_features=3, out_features=128, bias=True)
        (1): ReLU(inplace=True)
        (2): Linear(in_features=128, out_features=1024, bias=True)
        (3): ReLU(inplace=True)
        (4): Unflatten(dim=1, unflattened_size=(64, 4, 4))
    )
    (decoder): Sequential(
        (0): ConvTranspose2d(64, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
        (1): ReLU(inplace=True)
        (2): ConvTranspose2d(32, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
        (3): ReLU(inplace=True)
        (4): ConvTranspose2d(16, 1, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), output_padding=(1, 1))
        (5): Sigmoid()
    )
)

```

## C. Plot the latent space

```

In [84]: # Here are 10 colours you can use. But feel free to use others, too.
colour_options = ['k', 'tab:brown', 'r', 'orange', 'gold', 'lawngreen', 'forestgreen']

```

```

In [ ]: # Compute the latent-space representation for all the samples.

```

```

In [ ]: # You can create three 2D planar projections.

```

```

In [85]: # Or you can plot a 3D scatter plot.
fig = plt.figure(figsize=(7,6))
ax = fig.add_subplot(111, projection='3d')

===== YOUR CODE HERE =====

```

```

net.eval()
all_x, all_t = next(iter(train_all))
all_x = all_x.to(device)

with torch.no_grad():
    h = net.encode(all_x).cpu().numpy() # shape [1024, 3]
labels = all_t.numpy()

sc = ax.scatter(h[:,0], h[:,1], h[:,2],
                 c=labels, cmap='tab10', s=30, alpha=0.8)

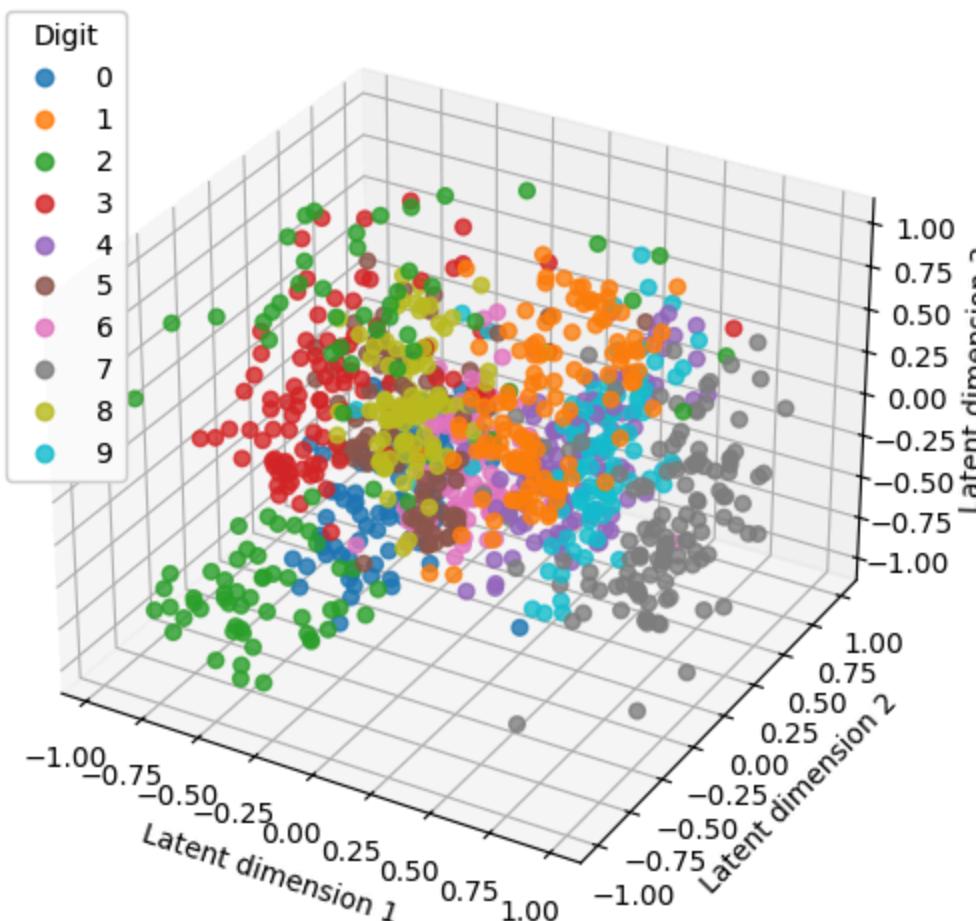
ax.set_xlabel('Latent dimension 1')
ax.set_ylabel('Latent dimension 2')
ax.set_zlabel('Latent dimension 3')
ax.set_title('3D latent representations of MNIST digits')

legend = ax.legend(*sc.legend_elements(), title="Digit")
ax.add_artist(legend)

plt.show()

```

3D latent representations of MNIST digits



## D. Plot reconstructed digit images

```
In [88]: net.eval()

    all_x, all_t = next(iter(train_all))
    all_x = all_x.to(device)
    all_t = all_t.numpy()

    with torch.no_grad():
        y = net(all_x).cpu()

    examples = []
    for digit in range(10):
        idx = np.where(all_t == digit)[0][0] # <-- fixed line
        examples.append((digit, all_x[idx], y[idx]))

    plt.figure(figsize=(10, 3))

    # Top row (inputs)
    for i, (digit, orig, recon) in enumerate(examples):
        plt.subplot(2, 10, i + 1)
        plt.imshow(orig.squeeze().cpu(), cmap='gray')
        plt.title(str(digit), color=colour_options[digit], fontsize=10)
        plt.axis('off')
        if i == 0:
            plt.ylabel("Input", fontsize=12)

    # Bottom row (reconstructions)
    for i, (digit, orig, recon) in enumerate(examples):
        plt.subplot(2, 10, 10 + i + 1)
        plt.imshow(recon.squeeze().cpu(), cmap='gray')
        plt.axis('off')
        if i == 0:
            plt.ylabel("Output", fontsize=12)

    plt.suptitle("MNIST digit reconstructions by class", fontsize=14)
    plt.tight_layout()
    plt.show()
```

MNIST digit reconstructions by class



Digit Reconstructions